The background of the slide is a dark, atmospheric aerial photograph of a river or stream flowing through a dense forest. The water is a deep, dark greenish-blue, and the surrounding trees are a mix of dark evergreens and lighter-colored deciduous trees, some with autumn foliage. The river curves from the bottom right towards the top left.

Relational Algebra, DDL, DML

CST 363

Database schema

- A **schema** is the design or structure of a specific database.
- An **instance** is a schema "instantiated" with data

Schema:

```
table roster {  
    class: string,  
    student: string  
}
```

Instance:

class	student
class_101	student_1
class_101	student_2
class_205	student_3
class_205	student_4
class_205	student_5

Database queries

- A query is a statement requesting the retrieval of information from a database.

Example SQL queries:

```
SELECT *
FROM customers
WHERE country = 'Guam';
```

```
SELECT *
FROM customers
WHERE city = 'Berlin' OR city = 'Munich';
```

Transactions

- A transaction is the *unit of change* in a database.
- Transactions contain a set of database operations.
- A transaction must succeed, or fail "completely".
- For example, a transaction might be used for moving a student from one class to another.

*more on transactions later in semester

A Relational DB is a group of tables

- A column is also called an **attribute**
- Each attribute has a **domain**, which is a set of values that are allowed in the column.
- A row is also called a **tuple**.

instructor

instructor_id	instructor_name	dept_name	salary
10101	Srinivasan	Comp. Sci.	null
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	null	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Tuples

Another table from the `courses` DB

- A tuple represents a relationship between the values of the tuple.
- A table represents a mathematical relation.
 - So a table is also called a "relation". This is why we say "relational database".

course

course_id	title	dept_name	credits
BIO-101	Intro. to Biology	Biology	4
BIO-301	Genetics	Biology	4
BIO-399	Computational Biology	Biology	3
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3
EE-181	Intro. to Digital Systems	Elec. Eng.	3
FIN-201	Investment Banking	Finance	3
HIS-351	World History	History	3
MU-199	Music Video Production	Music	3
PHY-101	Physical Principles	Physics	4

Null values

- Sometimes we don't know a value, or the value doesn't exist.
- In relational DBs we deal with this by using special `null` values.

instructor

instructor_id	instructor_name	dept_name	salary
10101	Srinivasan	Comp. Sci.	null
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	null	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Relation schema and instances

- **relation schema:** Gives the names and domains of the attributes (domains not shown here.)
 - `student(student_id, student_name, dept_name, tot_cred)`
- **relation instance:** an instance of a *relation schema* (also called a "table")

student

student_id	student_name	dept_name	tot_cred
128	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
19991	Brandt	History	80
23121	Chavez	Finance	110
44553	Peltier	Physics	56
45678	Levy	Physics	46

Database schema and instances

database schema: the relation schemas for all the relations in a database

```
student(student_id, student_name, dept_name, tot_cred)
department(dept_name, building, budget)
...
```

Database schema and instances (cont.)

database instance: all the tables of the database

student

student_id	student_name	dept_name	tot_cred
128	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
19991	Brandt	History	80
23121	Chavez	Finance	110
44553	Peltier	Physics	56
45678	Levy	Physics	46

department

dept_name	building	budget
Biology	Watson	90000
Comp. Sci.	Taylor	100000
Elec. Engr.	Taylor	85000
Finance	Painter	120000
History	Painter	50000
Music	Packard	80000
Physics	Watson	70000

All the names

In databases there seem to be two or three ways to say everything.

- **schema** = relation schema
- **table** = relation = relation instance = instance
- tuple = **row**
- **attribute** = column

We'll normally use the terms in bold

"Functions" on tables

- In programming we define objects, and methods on those objects.
- For example, for strings we have operations like
 - `substring(string, first, last)`
 - `concat(string1, string2)`
- In relational databases the objects are tables, and we define functions (operations) on tables.

"Selection" operator: Pick out rows of a table

patient

Patient No.	Last name	First name	Sex
454	Smith	John	M
223	Jones	Peter	M
597	Brown	Brenda	F
234	Jenkins	Alan	M
244	Wells	Christy	F

Sex = F

Patient No.	Last name	First name	Sex
597	Brown	Brenda	F
244	Wells	Christy	F

"Projection" operator: Pick out columns of a table

The operation takes a table as input and produces a table as output

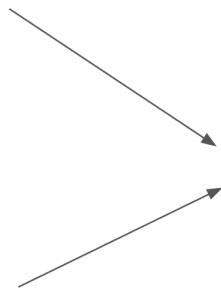
Patient No.	Last name	First name	Sex
454	Smith	John	M
223	Jones	Peter	M
597	Brown	Brenda	F
234	Jenkins	Alan	M
244	Wells	Christy	F

"Patient No.", "First name"

Patient No.	First name
454	John
223	Peter
597	Brenda
234	Alan
244	Christy

"Union" operator: Add the rows of two tables

Class_101	Student_1
Class_205	Student_4
Class_205	Student_5



Class_101	Student_1
Class_205	Student_4
Class_205	Student_5
Class_101	Student_2
Class_205	Student_3

Question about "projection" operator

Patient No.	Last name	First name	Ward No.
454	Smith	John	6
223	Jones	Peter	8
597	Brown	Brenda	3
234	Jenkins	Alan	7
244	Brown	Chris	6

What happens if we perform a `projection` on this table using attribute 'Last name'?

Last name
Smith
Jones
Brown
Jenkins
Brown

Is this a problem?

Removing duplicate rows

- As part of applying an operation of relation algebra, duplicate rows should be removed.
- Duplicate rows aren't technically allowed in a valid table.
- In practice a DB system may allow duplicate rows.

Defining the operations

```
selection(table, condition)
```

- return the rows of table that meet the condition
- example: `t = selection(student, tot_cred > 100)`

```
projection(table, attributes)
```

- returns the specified columns of table
- example: `t = projection(student, ["student_name", "student_id"])`

```
union(table1, table2)
```

- returns the table containing the rows of table1 and table2
- example: `t = union(student1, student2)`

Answering questions with the functions

What are the names of students with more than 100 credits?

- To get the answer:
 - select the rows of the `student` table with `tot_cred > 100`
 - get the `name` attribute

```
projection(selection(student, tot_cred > 100), ["student_name"])
```

Another example

What courses are offered in Spring '09, and what buildings are they in?

course_id	sec_id	semester	year	building	room_number	time_slot_id
BIO-101	1	Summer	2009	Painter	514	B
BIO-301	1	Summer	2010	Painter	514	A
CS-101	1	Fall	2009	Packard	101	H
CS-101	1	Spring	2010	Packard	101	F
CS-190	1	Spring	2009	Taylor	3128	E
CS-190	2	Spring	2009	Taylor	3128	A
CS-315	1	Spring	2010	Watson	120	D
CS-319	1	Spring	2010	Watson	100	B

```
t = selection(section, semester="Spring" and year=2009)
projection(t, ["course_id", "building"])
```

Exercise

Write the operations to get the IDs of students who took course CS-101 in Fall 2009.

ID	course_id	sec_id	semester	year	grade
128	CS-101	1	Fall	2009	A
128	CS-347	1	Fall	2009	A-
12345	CS-101	1	Fall	2009	C
12345	CS-190	2	Spring	2009	A
12345	CS-315	1	Spring	2010	A
12345	CS-347	1	Fall	2009	A

```
projection(selection(takes, course_id="CS-101"  
and semester="Fall" and year=2009), ["ID"])
```

A 'union' example

Get the ID of courses taught in Fall 2009 or Spring 2010

section

course_id	sec_id	semester	year	building	room_number	time_slot_id
BIO-101	1	Summer	2009	Painter	514	B
BIO-301	1	Summer	2010	Painter	514	A
CS-101	1	Fall	2009	Packard	101	H
CS-101	1	Spring	2010	Packard	101	F
CS-190	1	Spring	2009	Taylor	3128	E

```
t1 = selection(section, semester="Fall" and year=2009)
t2 = selection(section, semester="Spring" and year=2010)
projection(union(t1,t2), ["course_id"])
```

Last line could alternatively be written:

```
union(projection(t1,["course_id"]), projection(t2,["course_id"]))
```

Cartesian Product

id	name
1	Cesar
2	Rivka
3	Arturo

course	id	time
DB	1	10
OS	2	4

id	name	course	id	time
1	Cesar	DB	1	10
1	Cesar	OS	2	4
2	Rivka	DB	1	10
2	Rivka	OS	2	4
3	Arturo	DB	1	10
3	Arturo	OS	2	4

Question:

If you take the product of a table with 5 rows and a table with 8 rows, how many rows do you get?

Using cross product

What are the names of CS instructors and what classes do they teach?

instructor_id	instructor_name	dept_name	salary
10101	Srinivasan	Comp. Sci.	null
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	null	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-101	1	Spring	2010
15151	MUS-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

```
t1 = cross(instructor, teaches)
t2 = selection(t1, instructor.ID = teaches.ID and dept_name = "Comp. Sci.")
projection(t2, ["name", "course_id"])
```

Saying it in plain Greek

To make it look more like math, you can write the operations using Greek letters and symbols

- `selection(section, year=2015)`
 - $\sigma_{\text{year}=2015}(\text{section})$
 - sigma (σ) – "selection" starts with "s"
- `projection(section, ["course_id"])`
 - $\pi_{\text{course_id}}(\text{section})$
 - pi (π) – "projection" starts with "p"
- $\pi_{\text{course_id}}(\sigma_{\text{year}=2015}(\text{section}))$
 - First, filters the `section` table to only the rows where `year = 2015`.
 - Then, from those remaining rows, keep only the `course_id` column.

Part 2: SQL (DDL/DML)

What is SQL?

- SQL = "Structured Query Language"
- Originally developed at IBM in the early 70's
- Pronounced "sequel" or S-Q-L
- A special-purpose language that lets you express questions about your relational data in a precise way

SQL serves as the interface between the database and users/applications

Parts of SQL

SQL is not just for writing queries:

- Data-definition language (DDL)
 - create/delete/modify schemas
 - define integrity constraints
 - define views
 - drop tables
- Data-manipulation language (DML)
 - define queries
 - modify tables (insert/delete/modify tuples)
- TCL (Transaction Control Language): BEGIN , COMMIT , ROLLBACK
- DCL (Data Control Language): GRANT , REVOKE

Defining a relation schema

Question: What information do you provide when you define a relation schema?

- name of the table
- names and domains of each attribute
- primary key
- optionally:
 - whether an attribute is allowed to have a null value
 - foreign keys

SQL for defining a relation schema

Relation schema: `department(dept_name, building, budget)`

DDL table definition:

```
CREATE TABLE department (
    dept_name VARCHAR(20) PRIMARY KEY,
    building VARCHAR(15),
    budget NUMERIC(12, 2)
);
```

- For each attribute we give the attribute name and type.
- We also give the primary key.

Another 'create table' example

```
CREATE TABLE student (
    student_id VARCHAR(20) PRIMARY KEY,
    name VARCHAR(20) NOT NULL,
    dept_name VARCHAR(20),
    tot_cred NUMERIC(3, 0),
    FOREIGN KEY (dept_name) REFERENCES department(dept_name)
);
```

- The foreign key constraint says:
 - every `dept_name` value in the student table must equal the key of some row in the department table.
- The `NOT NULL` constraint says:
 - no name value in the student table can be `NULL`
- PostgreSQL requires referenced tables exist before creating FKS – `department` needs to be created first

SQL Constraints Review

- PRIMARY KEY = UNIQUE + NOT NULL
- FOREIGN KEY enforces referential integrity
- NOT NULL disallows missing values

Attribute types

Recall that attributes can only be simple values, like integers, and strings.*

Here are some of the main types allowed by SQL:

- `VARCHAR(n)` : a variable-length string of length at most `n`
- `INT` : an integer – `integer` is also okay
- `NUMERIC(p, d)` : a fixed-point number of `p` total digits, with `d` digits to right of the decimal point.
 - Example: 123.45 is `numeric(5,2)`
- `FLOAT(p)` : a floating-point number with either 4 byte or 8 byte precision
 - alternatively can use `REAL` (4 bytes) and `DOUBLE PRECISION` (8 bytes)

*In the relational model, attributes are atomic (simple) values.

PostgreSQL also supports richer types (e.g., arrays/JSON), but we'll focus on relational design first.

Yet another example

```
CREATE TABLE takes (
    student_id VARCHAR(5),
    course_id VARCHAR(8),
    section_id VARCHAR(8),
    semester VARCHAR(6),
    section_year NUMERIC(4, 0),
    grade VARCHAR(2),
    PRIMARY KEY (student_id, course_id, section_id, semester, section_year),
    FOREIGN KEY (course_id, section_id, semester, section_year)
        REFERENCES section(course_id, section_id, semester, section_year) ON DELETE CASCADE,
    FOREIGN KEY (student_id)
        REFERENCES student(student_id) ON DELETE CASCADE
);
```

- primary key has multiple attributes
- two foreign keys

Cascaded deletes

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-101	1	Spring	2010
15151	MUS-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

- `ID` in the **teaches** table is a foreign key that "points to" the `ID` field of the **instructor** table.
- What to do when the first row is deleted from the instructor table?
 - Default: delete is rejected if referenced.
 - **CASCADE** : referencing rows are deleted automatically.
 - **CASCADE** is powerful; use intentionally.

Inserting rows into a table

patient_no	last_name	first_name	sex	date_of_birth	ward
454	Smith	John	M	14.08.78	6
223	Jones	Peter	M	07.12.85	8
597	Brown	Brenda	F	17.06.61	3
234	Jenkins	Alan	M	29.01.72	7
244	Wells	Chris	F	25.02.95	6

```
CREATE TABLE patient (
    patient_no INTEGER PRIMARY KEY,
    last_name VARCHAR(64) NOT NULL,
    first_name VARCHAR(64) NOT NULL,
    sex VARCHAR(1),
    date_of_birth DATE,
    ward INTEGER
);
```

Inserting rows into a table (cont.)

SQL also has statements to delete and modify rows, as well as batch operations.

```
CREATE TABLE patient (
    patient_no INTEGER PRIMARY KEY,
    last_name VARCHAR(64) NOT NULL,
    first_name VARCHAR(64) NOT NULL,
    sex VARCHAR(1),
    date_of_birth DATE,
    ward INTEGER
);
```

```
INSERT INTO patient VALUES (454, 'Smith', 'John', 'M', '1978-08-14', 6);
INSERT INTO patient VALUES (223, 'Jones', 'Peter', 'M', '1985-12-07', 8);
INSERT INTO patient VALUES (597, 'Brown', 'Brenda', 'F', '1961-06-17', 3);
INSERT INTO patient VALUES (234, 'Jenkins', 'Alan', 'M', '1972-01-29', 7);
INSERT INTO patient VALUES (244, 'Wells', 'Chris', 'F', '1995-02-25', 6);
```

More efficient insert

- PostgreSQL supports multi-row insert

```
INSERT INTO patient (patient_no, last_name, first_name, sex, date_of_birth, ward)
VALUES
(454, 'Smith', 'John', 'M', '1978-08-14', 6),
(223, 'Jones', 'Peter', 'M', '1985-12-07', 8);
```

Deleting or Modifying a table

```
DROP TABLE IF EXISTS instructor;
```

- delete the `instructor` schema and its tuples

```
ALTER TABLE instructor ADD office VARCHAR(5);
```

- adds attribute 'office' to the instructor schema
- value of office is initialized to 'null'

Summary

- `CREATE TABLE` – define a relation schema
- `INSERT` – insert rows into a relation instance
- `DROP TABLE` – delete relation schemas
- `ALTER TABLE` – modify relation schemas